

Exercise 1

Jay Morgan

October 18, 2018

R Markdown

Task 1 Sequences

1. Create the vector `x`.

```
x <- seq(-30, 30, 5)
```

```
x
```

```
## [1] -30 -25 -20 -15 -10 -5 0 5 10 15 20 25 30
```

2. Create a second vectory, which contains exactly the same elements with the same multiplicities as vector `x`, but in a *random* order.

```
y <- seq(-30, 30, 5)
```

```
set.seed(283)
```

```
y <- sample(y)
```

```
y
```

```
## [1] 20 -25 0 5 -15 10 -10 -30 -5 15 30 -20 25
```

3. On how many and which position(s) do the values of `x` and `y` agree?

```
length(which(x == y))
```

```
## [1] 2
```

```
which(x == y)
```

```
## [1] 2 10
```

Task 2 Sequences

1. John Wallis (1616-1703):

```
i <- 1:10000

Wallis <- prod(((2*i)/((2*i)-1))*((2*i)/((2*i)+1)))
Wallis

## [1] 1.570757
```

2. Gottfried Leibnitz (1646-1716):

```
i <- 1:10000

Leibnitz <- sum((-1)^(i+1)/((2*i)-1))
Leibnitz

## [1] 0.7853732
```

3. Leonhard Euler (1707-1783):

```
i <- 1:10000

Euler <- sum(1/(i^2))
Euler

## [1] 1.644834
```

Which of the above formulas shows the smallest relative deviation from π for this value of n ?

```
pi

## [1] 3.141593

Wallis*2

## [1] 3.141514

Leibnitz*4

## [1] 3.141493

sqrt(Euler*6)

## [1] 3.141497
```

The Wallis estimation has the smallest deviation from π at $n = 10000$.

Task 3 Vectors

1. Create a data vector x containing the natural numbers $1, 2, \dots, 100$ and a second data vector y containing a sample of size $n = 70$ from the set of natural numbers $1, 2, \dots, 150$ with replacement.

```
x <- 1:100

set.seed(283)
y <- sample(1:150, 70, replace = TRUE)
```

2. Determine those elements of x , which are not contained in y . How many elements are there?

```
setdiff(x, y)

## [1]  2  3  4  7  8  9 10 11 12 14 15 19 21 22 23 26 27
## [18] 30 31 32 33 35 36 39 40 42 44 45 47 48 49 51 52 53
## [35] 55 57 58 59 62 63 64 66 67 68 69 70 71 73 74 75 76
## [52] 77 78 79 80 81 82 83 84 85 89 93 94 95 96 97 98 100

length(setdiff(x, y))

## [1] 68
```

3. Are there duplicate entries in your vector y ? If yes, create a new vector z , with these duplicates. If no, just take y as vector z .

```
any(duplicated(y))

## [1] TRUE

z <- which(duplicated(y) == 'TRUE')

z

## [1] 14 20 22 23 25 27 28 34 41 47 52 55 58 59 60 61 67
```

4. How many elements of z are multiples of 3?

```
length(which(z%3 == 0))

## [1] 2
```

5. Revert the vector y without (!) using the function `rev()`, i.e. if y were the vector $(5, 20, 81)$, the result should be the vector $(81, 20, 5)$.

```
y[seq(70, 1, -1)]

## [1] 122 65 13 88 92 60 111 20 46 109 34 133 54 6 105 117 124
## [18] 109 120 41 136 142 107 24 34 5 29 72 37 17 54 17 123 91
## [35] 99 117 102 18 133 50 43 102 28 127 1 25 86 127 147 88 28
## [52] 145 138 134 90 132 61 147 28 38 139 24 25 104 61 56 120 87
## [69] 16 127
```

Task 4 Point Estimation

1. Draw a reproducible sample of size $n=30$ from a normal distribution with $\mu=7$ and $\sigma^2=4$.

```
n <- 30
set.seed(283)
rnorm(n, mean = 7, sd = 4)

## [1] 11.0274794  7.8015545  5.6970847  8.9653964  2.9961200  4.3336843
## [7] 15.1438518 11.6982887 11.8562358 14.0455312  7.8726079 11.0291938
## [13]  3.1068399 10.9842637  8.8247529  5.2080525  2.2740651 10.0716705
## [19]  8.0364105  2.0504061  2.0327303  6.7660703 -0.3877495  2.9717444
## [25] 13.2750485  4.5706240  9.3318184 10.0766969 -0.1899235 11.7089006
```

2. Estimate μ and σ^2 on the basis of your sample using the above formulas, i.e. without (!) using the functions `mean()`, `var()` and `sd()`.

```
n <- 30
set.seed(283)
x <- rnorm(n, mean = 7, sd = 4)

sum(x)/n      # Estimated mean

## [1] 7.439315

sum((x-sum(x)/n)^2)/(n-1) # Estimated variance

## [1] 18.72024
```

3. Compare your results with the output of the functions `mean()` and `var()`.

```
n <- 30
set.seed(283)
x <- rnorm(n, mean = 7, sd = 4)

sum(x)/n      # Estimated mean

## [1] 7.439315

mean(x)       # Population mean

## [1] 7.439315

sum(x)/n == mean(x) # Comparison of mean

## [1] FALSE

sum((x-sum(x)/n)^2)/(n-1) # Estimated variance

## [1] 18.72024

var(x)        # Population variance

## [1] 18.72024

sum((x-sum(x)/n)^2)/(n-1) == var(x) # Comparison of variance

## [1] TRUE
```

Task 4 continued

4. Are your estimates close to the population values? Repeat the steps 1 and 3 from above with a sample of size $n = 3000$. What do we learn?

The estimates are very close to the population values produced by the functions.

```
n <- 3000
set.seed(283)
x <- rnorm(n, mean = 7, sd = 4)

sum(x)/n    # Estimated mean

## [1] 6.969209

mean(x)     # Population mean

## [1] 6.969209

sum(x)/n == mean(x)  # Comparison of mean

## [1] TRUE

sum((x-sum(x)/n)^2)/(n-1)  # Estimated variance

## [1] 16.12216

var(x)      # Population variance

## [1] 16.12216

sum((x-sum(x)/n)^2)/(n-1) == var(x)  # Comparison of variance

## [1] TRUE
```

We learn that as the number of samples increase, the observed mean and variance values become closer to the defined μ and σ^2 values of 7 and 4 respectively.

Task 5 Interval Estimation

1. Draw a reproducible sample of size $n=30$ from a normal distribution with $\mu=7$ and $\sigma^2=4$.

```
n <- 30
set.seed(283)
rnorm(n, mean = 7, sd = 4)

## [1] 11.0274794  7.8015545  5.6970847  8.9653964  2.9961200  4.3336843
## [7] 15.1438518 11.6982887 11.8562358 14.0455312  7.8726079 11.0291938
## [13]  3.1068399 10.9842637  8.8247529  5.2080525  2.2740651 10.0716705
## [19]  8.0364105  2.0504061  2.0327303  6.7660703 -0.3877495  2.9717444
## [25] 13.2750485  4.5706240  9.3318184 10.0766969 -0.1899235 11.7089006
```

2. Calculate a confidence interval for μ and σ^2 for $\alpha = 0.05$ (hence the confidence level is $1 - \alpha = 0.95$). R-functions such as `mean()`, `sd()` and `var()` are allowed.

```
n <- 30
set.seed(283)
x <- rnorm(n, mean = 7, sd = 4)

qnorm(p = c(0.05, 0.95), mean = 7, sd = 4)

## [1]  0.4205855 13.5794145

qnorm(p = c(0.05, 0.95), mean = mean(x), sd = sd(x))

## [1]  0.3225442 14.5560858
```

3. Do the true parameters lie in your confidence intervals? If yes, is this always the case? If no, why not?

```
mean_95 <- qnorm(p = c(0.05, 0.95), mean = 7, sd = 4)

mean_x <- qnorm(p = c(0.05, 0.95), mean = mean(x), sd = sd(x))

mean_95[1] < mean_x[1]

## [1] FALSE

mean_95[2] > mean_x[2]

## [1] FALSE
```